

```
#Rishikesh_2412res99

# Import libraries
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

# Load the data from CSV file
data = pd.read_csv("/content/HR_comma_sep.csv")

# Quick look at the data
print(data.head())

    satisfaction_level  last_evaluation  number_project  average_montly_hours \
0            0.38           0.53             2                  157
1            0.80           0.86             5                  262
2            0.11           0.88             7                  272
3            0.72           0.87             5                  223
4            0.37           0.52             2                  159

   time_spend_company  Work_accident  left  promotion_last_5years Department \
0                   3           0      1                  0        sales
1                   6           0      1                  0        sales
2                   4           0      1                  0        sales
3                   5           0      1                  0        sales
4                   3           0      1                  0        sales

   salary
0   low
1 medium
2 medium
3   low
4   low

# Split data into employees who left (1) and stayed (0)
left = data[data["left"] == 1]
stayed = data[data["left"] == 0]

# Calculate means for key variables
print("Employees who LEFT:")
print(left[["satisfaction_level", "last_evaluation", "number_project",
           "average_montly_hours", "time_spend_company", "Work_accident",
           "promotion_last_5years"]].mean())

print("\nEmployees who STAYED:")
print(stayed[["satisfaction_level", "last_evaluation", "number_project",
              "average_montly_hours", "time_spend_company", "Work_accident",
              "promotion_last_5years"]].mean())

# Observations (you can read these in output):
# - Low satisfaction, extreme projects/hours, mid-tenure, no promotions linked to leaving

    Employees who LEFT:
satisfaction_level      0.440098
last_evaluation         0.718113
number_project          3.855503
average_montly_hours    207.419210
time_spend_company       3.876505
Work_accident           0.047326
promotion_last_5years   0.005321
dtype: float64

Employees who STAYED:
satisfaction_level      0.666810
last_evaluation          0.715473
number_project           3.786664
average_montly_hours     199.000203
time_spend_company        3.380032
Work_accident             0.175009
promotion_last_5years    0.026251
dtype: float64

# Group by salary and calculate proportion who left
salary_retention = data.groupby("salary")["left"].mean()

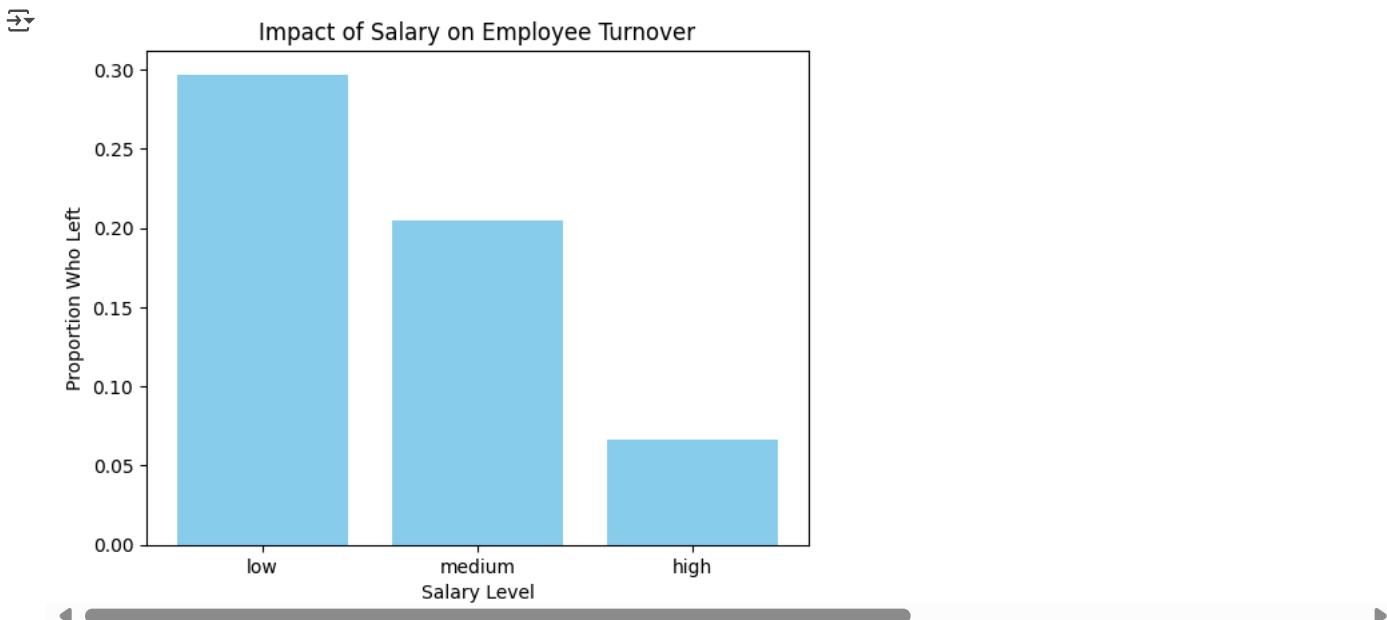
# Order salary levels: low, medium, high
salary_retention = salary_retention.reindex(["low", "medium", "high"])

# Plot bar chart
plt.bar(salary_retention.index, salary_retention.values, color="skyblue")
```

```

plt.title("Impact of Salary on Employee Turnover")
plt.xlabel("Salary Level")
plt.ylabel("Proportion Who Left")
plt.show()

```

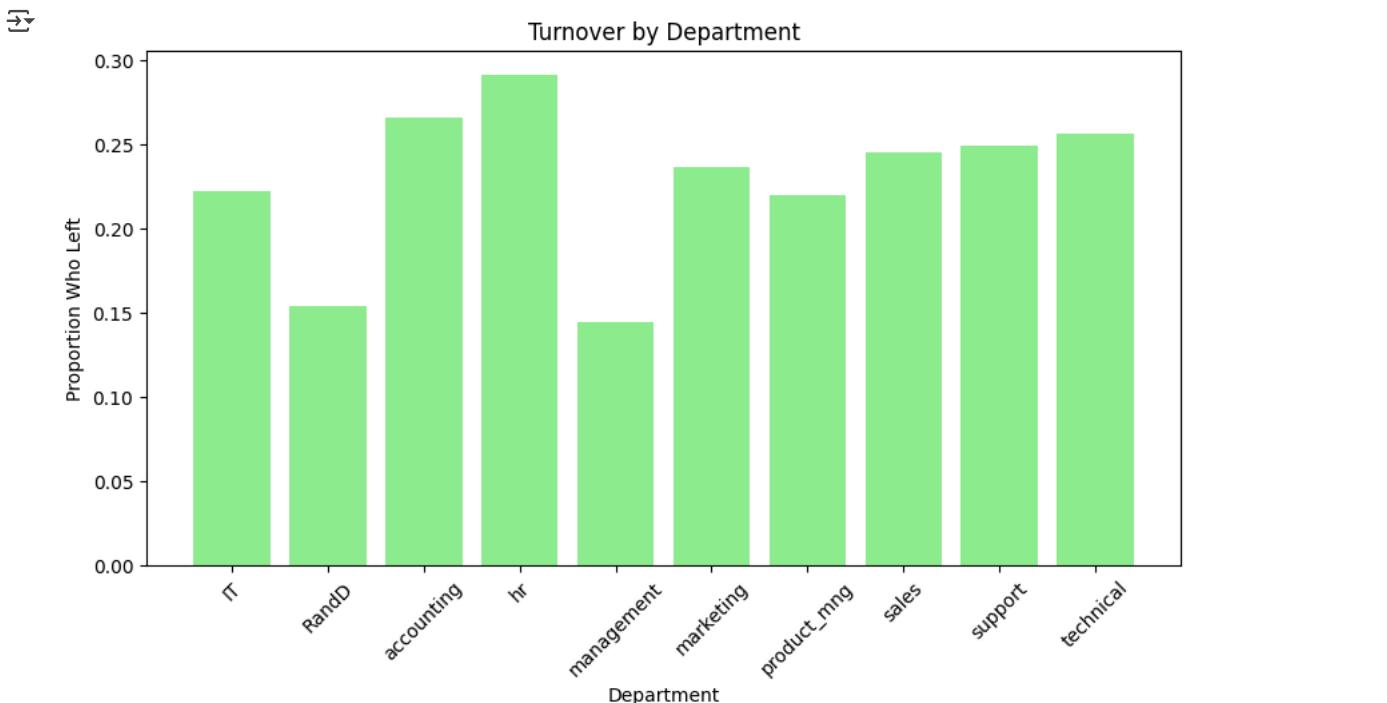


```

# Group by department and calculate proportion who left
dept_retention = data.groupby("Department")["left"].mean()

# Plot bar chart
plt.figure(figsize=(10, 5)) # Make it wider to fit all departments
plt.bar(dept_retention.index, dept_retention.values, color="lightgreen")
plt.title("Turnover by Department")
plt.xlabel("Department")
plt.ylabel("Proportion Who Left")
plt.xticks(rotation=45) # Rotate labels for readability
plt.show()

```



```

# Prepare the data
# Convert salary to numbers (low=0, medium=1, high=2)
data["salary_num"] = data["salary"].map({"low": 0, "medium": 1, "high": 2})

# Select features (X) and target (y)
features = ["satisfaction_level", "number_project", "average_montly_hours",
            "time_spend_company", "promotion_last_5years", "salary_num"]
X = data[features]
y = data["left"]

# Split data into training (70%) and testing (30%)

```

```
# Split data into training (70%) and testing (30%)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

# Create and train the model
model = LogisticRegression(max_iter=1000) # max_iter avoids convergence warning
model.fit(X_train, y_train)

print("Model trained successfully!")

→ Model trained successfully!

# Make predictions on test data
y_pred = model.predict(X_test)

# Calculate accuracy
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f} ({accuracy * 100:.1f}%)")

→ Model Accuracy: 0.77 (76.6%)
```